

Dipayan Bhowmick

Head of Engineering | AI, Data & Agentic Systems

Bangalore, India | +91-9731099002 | me@dipayanb.com | [LinkedIn](#) | [GitHub](#)

Summary

Engineering leader who builds AI and data products from zero to revenue. Joined Acceldata as employee #2 and grew the Data Observability platform to \$4M ARR across 8+ Fortune 500 accounts; most recently led engineering and the India site for Altimate AI, shipping agentic AI teammates for data teams. 18 years across the stack, currently hands-on with LLMs, RAG, and agent frameworks, with deep roots in distributed data systems. Comfortable owning architecture, team, and GTM in equal measure.

Work Experience

Head of Engineering & Site Lead (India) | Altimate AI

Feb 2025 – Jan 2026 (1 yr) | Bengaluru

- Led engineering for a seed-stage AI startup building agentic AI teammates for data teams, integrated into IDEs, Git, and Slack; shipped to 5+ customers in production
- Designed and built the memory system underpinning the AI agents, the layer that lets them carry context across sessions, tools, and data sources
- Built the SQL parser and scanner for the DataPilot platform, enabling the agents to reason over customer data warehouses safely and accurately
- Set up the India site from scratch and built a team of 20 across engineering, product, and operations; managed vendors, hiring, and site strategy end-to-end
- Introduced engineering processes calibrated to the company's stage, enough rigor to ship reliably, light enough not to slow a 20-person team
- Migrated ingestion to ClickHouse, cutting OpenSearch usage to 10% and improving pipeline reliability; reduced AWS costs by 25% through resource auditing

Skills: Python, FastAPI, LangChain, CrewAI, ClickHouse, OpenSearch, AWS, AI Agents, RAG Pipelines, Tools/MCPs, SQL Parsing, People Management, Vendor Management, Process Engineering, Customer Success

Director of Engineering, Intel Exchange | Cyware

Sep 2024 – Feb 2025 (6 mos) | Bengaluru

- Led engineering for the Intel Exchange threat intelligence platform with a team of 25 across backend, frontend, and QA, partnering with PMs and GTM to ship customer-driven roadmap items
- Drove the platform to FedRAMP certification, the credential required to unlock a major US government opportunity
- Delivered the 2.x → 3.x upgrade path, unblocking customer migrations that had previously had no clean route forward
- Built out platform observability from scratch, cutting time-to-diagnose for production issues and reducing escalations to engineering
- Shipped major releases, platform stabilization, and architectural improvements alongside FedRAMP work, keeping the existing customer base on a steady release cadence while standing up the compliance posture for new ones

Skills: Python, FastAPI, Django, Cybersecurity, Threat Intelligence, Observability, Engineering Processes, GTM Alignment

Startup Attempt (Did Not Continue) | AI Revenue Operations

Aug 2023 – Sep 2024 (1 yr 2 mos) | Bengaluru

- Built an AI-native revenue operations product to a working POC with a co-founder and one design partner; pursued fundraising in parallel, and wound down when the signal didn't justify continuing

Director of Engineering (Principal Engineer → Senior EM → Director) | Acceldata

Aug 2019 – Aug 2023 (4 yrs) | Bangalore

Joined as the second employee and built the Data Observability Cloud Platform (ADOC) from inception to \$4M+ ARR.

- Pioneered and architected the ADOC platform, owning roadmap and engineering development end-to-end
- Partnered across Sales, Marketing, and Customer Success to win and scale more than 8 Fortune 500 accounts
- Scaled from a two-person team to 20 engineers across three teams; recruited 75+ people across cross-functional roles in the organization
- Architected the cloud SaaS version of the product and scaled that to 20+ customers
- Reduced the product's release cycle from once a quarter to once in 2 weeks
- Engineered the Alerts and auto-actions framework, the system sends more than 50,000 emails and Slack messages every day
- Ran Spark on Kubernetes before Spark had native K8s support
- For one customer, processed 450 million nested JSON records per batch and reduced processing time from 20 days to 7 hours

Skills: Apache Spark, Kotlin, Scala, Postgres, MongoDB, Snowflake, Databricks, Kafka, Microservices, Kubernetes, CI/CD, Machine Learning, Growth Strategies, Budgeting, People Management, Presales

Engineering Manager (Senior MTS → Staff Engineer → EM) | Hortonworks

Jul 2015 – Mar 2019 (3 yrs 9 mos) | Bangalore

Grew from individual contributor on Apache Ambari to managing a team of 8 engineers building Dataplane.

- Scaled Apache Ambari to manage 4,000-node Hadoop clusters; recognized with Apache committership
- Led the development of the Ambari Views framework, improving performance and user experience across Hive and Files views
- Built Dataplane, a tool for orchestrating software deployments within the Hadoop ecosystem; deployed to three customer datacenters
- Led the effort to make Dataplane a closed-source product within an open-source company

Skills: Java, Scala, Kafka, Docker, Kubernetes, Postgres, Hadoop, Apache Ambari

Earlier Experience

- **SDE III | Flipkart** (Dec 2014 – Jul 2015): Managed five order management microservices handling 500,000 orders/day at peak
- **Team Lead | [24]7.ai (ILabs)** (Jun 2013 – Dec 2014): Built data analysis tooling for the data science team and a customer-facing reporting platform
- **Senior Member Technical Staff | Oracle** (Sep 2011 – Jun 2013)
- **Various Roles | GE Healthcare, Mobilium (Roamware), IBM** (Jul 2007 – Sep 2011)

Education

National Institute of Technology Jamshedpur

B.Tech (Hons), Computer Science Engineering | 2003 – 2007 | Member of the Programming Society

Skills

AI & Agentic Systems

AI Agents, LLMs, RAG Pipelines, Tools/MCPs, Machine Learning

Data Infrastructure & Observability

Data Observability, Data Pipelines, Apache Spark, ClickHouse, Kafka, Snowflake, Databricks, Hadoop Ecosystem

Cloud, DevOps & Architecture

AWS, Azure, Kubernetes, Docker, CI/CD, Distributed Systems, Microservices

Languages & Databases

Python, Kotlin, Scala, Java, Typescript/JavaScript, Postgres, MySQL, MongoDB, Redis

Management

People Management, Hiring, Growth Strategies, Budgeting, Pre-sales, Customer Success, Enterprise Architecture

Personal Projects

- **AI Voice Receptionist:** A naive voice receptionist built on ElevenLabs and Cartesia APIs, using Elixir and WebSockets
 - **MeldUI:** A Vue.js component library based on shadcn, TanStack Tables, and ECharts
 - **Personal Data Analyst:** Local CSV/JSON analysis using natural language, built with Tauri (Rust), Vue.js, DuckDB, and Anthropic
 - **DiningBuddy.ca:** Built the AI interface and data ingestion pipeline for a restaurant discovery platform
 - **Open Source:** Contributed bug reports and PRs to [Jido](#) (Elixir AI agent framework) and [Kreuzberg](#) (document text extraction)
-

Certifications & Affiliations

- **Apache Ambari Committer:** The Apache Software Foundation (Jan 2016 – Present)
- **Generative AI with Large Language Models:** DeepLearning.AI & Amazon Web Services (Mar 2024)